

WEBVTT

00:01.600 --> 00:05.720

In this video, we examine the component architecture of QRadar SIEM.

00:06.960 --> 00:11.040

After looking at the high-level architecture, we focus on the details of the Flow

00:11.080 --> 00:14.080

and Event Collectors, the Event Processor, and the Console.

00:15.840 --> 00:18.320

Let's begin by looking at the high-level architecture.

00:19.840 --> 00:24.920

Events from individual log sources and network flow data are collected by the QRadar Event

00:24.960 --> 00:28.000

and Flow Collectors, including QRadar Network Insights.

00:29.360 --> 00:33.200

Once the flow and event data is forwarded to the Event or Flow Processor,

00:33.360 --> 00:36.360

it is stored in the Ariel database on that same Processor.

00:36.720 --> 00:40.760

To fulfill the tamper-proof data storage aspects for compliance mandates,

00:41.560 --> 00:46.360

data cannot be changed as soon as it is stored in the Ariel flow or event database tables.

00:47.520 --> 00:52.880

If accumulation is required, the accumulated data is stored in Ariel accumulation data tables.

00:53.880 --> 00:55.560

At any point in time,

00:55.600 --> 01:00.200

data can be selectively indexed to support specific search and report requirements.

01:01.200 --> 01:07.160

If needed, one or multiple Data Nodes can be attached to the Event Processor to extend the storage capacity.

01:08.760 --> 01:12.000

Once the event and Flow Processor has finished processing the data,

01:12.240 --> 01:15.920

the QRadar Console services continue with consolidated processing.

01:16.480 --> 01:19.560

Here, the Magistrate component creates and manages offenses.

01:20.480 --> 01:24.600

These offenses, as well as asset, identity, and configuration information,

01:24.840 --> 01:28.200

are stored in the Master PostgreSQL database on the Console.

01:29.000 --> 01:35.440

In every QRadar SIEM deployment, there is one master PostgreSQL database with optional read-only copies

01:35.480 --> 01:40.640

on each processor for backup and automatic restore (these are not depicted in this diagram).

01:41.400 --> 01:45.960

Apps from the IBM Security App Exchange can be hosted and run on the QRadar Console.

01:46.720 --> 01:50.040

In order to free up capacity and improve performance on the Console,

01:50.240 --> 01:53.920

apps can be hosted on a dedicated App Host appliance.

01:53.960 --> 01:59.120

Secure SSH communication between appliances in a distributed environment is supported.

01:59.840 --> 02:02.160

Let's now look at the details of the Flow Collector.

02:03.640 --> 02:10.200

A network flow record provides information about a conversation between two devices by using a specific protocol,

02:10.400 --> 02:13.680  
and it can include fields  
that provide details about the conversation.

02:14.680 --> 02:18.960  
Examples include the source and destination  
IP addresses, ports, and other fields.

02:19.960 --> 02:25.880  
Flow data packets can be collected from a variety of network device  
vendors and directly from the network interface.

02:26.920 --> 02:32.440  
Collected flow data can update asset profiles  
with the ports and services that are running on each host.

02:32.480 --> 02:40.040  
If a new host is detected through network flow data,  
a new asset is created in the QRadar asset database.

02:40.280 --> 02:42.480  
Next in line is the Aggregator.

02:42.520 --> 02:48.480  
This component enforces the license limit for the Flow Collector,  
which is measured in Flows Per Minute (FPM).

02:49.440 --> 02:54.040  
If the license limit is exceeded, flows are  
temporarily stored in an overflow buffer,

02:54.040 --> 02:56.320  
which processes them with the next set of flows.

02:56.960 --> 03:04.480  
Additionally, an overflow record is created  
with the source and destination address 127.0.0.4/5.

03:05.120 --> 03:08.760  
Every flow source  
protocol has an overflow buffer of 5 GB,

03:08.960 --> 03:12.520  
and if the overflow buffer fills up,  
the additional flows are dropped.

03:14.400 --> 03:18.640  
More information about the QRadar licenses and handling the overflow  
buffer

03:18.640 --> 03:22.040  
was presented in the previous QRadar SIEM Architecture video.

03:23.160 --> 03:28.920

Then the Application Detection Module uses four methods to determine the application based on the flow data.

03:30.080 --> 03:32.480

The first is the User Defined method.

03:32.520 --> 03:37.160

This method is used when users have a proprietary application that runs on their network.

03:37.720 --> 03:46.880

For example, all traffic that goes to host 10.100.100.42 on port 443 is recognized to be MySpecialApplication.

03:47.880 --> 03:50.600

The second method uses State-based decoders.

03:51.120 --> 03:53.920

This method is implemented by looking at the source code.

03:54.600 --> 03:58.520

It determines the application by analyzing the payload for multiple markers.

03:59.600 --> 04:08.480

For example, if you see A followed by B, then application = X; and if you see A followed by C, then application = Y.

04:09.240 --> 04:11.240

The next method uses Signature matching.

04:11.760 --> 04:14.840

This method relies on basic string matching in the payload.

04:15.400 --> 04:19.840

See the Applications configuration guide for defining application mappings and signatures.

04:21.000 --> 04:23.480

The final method uses Port-based matching.

04:23.840 --> 04:29.960

In this case, applications are matched based on their port use, for example, port 80 = HTTP.

04:31.280 --> 04:35.160

Finally, the flow data packets reach the Flow reporting and routing component.

04:36.320 --> 04:39.400

This component is responsible for creating superflows.

04:39.440 --> 04:44.240  
Superflows only store a single flow record  
with the collection of IP addresses,

04:44.280 --> 04:48.640  
which allows processing of flows  
to be faster and require less storage space.

04:49.600 --> 04:51.840  
There are three types of superflows:

04:51.840 --> 04:58.200  
Type A superflows contain a single source and multiple destination  
addresses with the same destination port,

04:58.200 --> 05:00.920  
byte count, and source flags or ICMP codes.

05:01.520 --> 05:04.400  
An example for a type  
A superflow is a network sweep.

05:04.920 --> 05:09.440  
Type B superflows contain multiple sources  
and a single destination address

05:09.680 --> 05:14.360  
with the same destination port,  
byte count, and source flags for ICMP codes.

05:15.480 --> 05:19.560  
An example for a type B superflow is a  
Distributed Denial of Service (DDoS) attack.

05:21.320 --> 05:23.840  
And Type C superflows contain a single source

05:24.080 --> 05:27.560  
and destination address  
with changing source and destination ports.

05:28.200 --> 05:31.760  
An example for a type C superflow is a port scan.

05:31.800 --> 05:37.520  
Specific rule tests can leverage the flow type  
to determine whether an offense needs to be created.

05:38.160 --> 05:40.280  
The creation of superflows can be disabled.

05:41.480 --> 05:45.240  
Up to a configurable number of bytes,

if the payload is unencrypted,

05:45.280 --> 05:49.040

QRadar QFlow provides

Layer 7 insights into the payload.

05:49.560 --> 05:55.880

By using a Tap or SPAN port, QFlow collects raw packets and places them into 60-second chunks.

05:56.920 --> 06:07.920

QFlow can also receive Layer 4 flows from other network devices in IPFIX/Netflow, sFlow, J-Flow Packeteer, and Flowlog file accounting technologies.

06:09.240 --> 06:11.120

Let's now look at the Event Collector.

06:13.040 --> 06:16.920

Each Event Collector gathers events from local and remote log sources.

06:17.720 --> 06:21.680

Once the raw data packets are received,  
the license limit is checked first.

06:22.280 --> 06:26.640

On the Event Collector,  
this limit is measured in events per second (EPS).

06:27.120 --> 06:33.880

If the EPS license is exceeded, they are temporarily stored in an overflow buffer and then processed during the next cycle.

06:34.400 --> 06:36.360

If the overflow buffer fills up,

06:36.400 --> 06:40.280

the additional events are dropped  
and a message is logged for the administrators.

06:40.800 --> 06:44.600

More information about the QRadar licenses  
and handling of the overflow

06:44.640 --> 06:48.400

buffer was presented in the previous QRadar  
Architecture video.

06:49.680 --> 06:54.840

Log sources are automatically discovered after  
record analysis in the Traffic Analysis module.

06:55.600 --> 07:00.200

This is an essential module

for automating a successful evaluation or deployment

07:00.400 --> 07:04.760

because it categorizes traffic from devices  
that are unknown to the system.

07:05.600 --> 07:11.400

If detection is successful on an IP address, log  
source detection creates a new QRadar log source.

07:12.640 --> 07:16.760

The Traffic Analysis module only carries out  
detection on event protocols

07:16.800 --> 07:20.960

that are pushed to the Event Collector,  
for example, syslog.

07:21.000 --> 07:25.360

After the correct log source is detected,  
such as a Checkpoint firewall,

07:25.560 --> 07:28.440

the individual Device Support Modules parse  
the events.

07:29.120 --> 07:32.120

First, the events are normalized  
where source specific

07:32.160 --> 07:36.160

data fields are mapped into QRadar terminology  
for further processing.

07:36.760 --> 07:44.200

The log source parser then extracts the log source event ID from the log  
record and maps that to the QRadar identifier (QID).

07:44.880 --> 07:49.280

This is a unique ID that links the extracted  
log source event ID to a QID.

07:50.400 --> 07:56.480

Each QID relates to a custom event name and description,  
as well as severity and event category information.

07:56.920 --> 08:02.000

The event category information is structured into High Level Categories  
(HLC) and Low Level Categories (LLC).

08:02.680 --> 08:07.760

For example, a valid category combination  
is "Authentication" (being a High Level Category)

08:07.960 --> 08:11.120  
and "Admin Login Successful"  
(being a Low Level Category).

08:11.960 --> 08:17.800  
Finally, the coalescing filter can optionally bundle identical events to conserve system usage

08:17.800 --> 08:20.320  
before handing off the data to the Event Processor.

08:22.760 --> 08:24.680  
The Event Processor can receive

08:24.720 --> 08:29.520  
event and flow data from Event and Flow Collectors  
as well as other Event Processors

08:29.560 --> 08:33.120  
that might be distributed  
throughout the organization's IT deployment.

08:33.640 --> 08:38.560  
First, the Overflow Filter enforces the license  
in a similar way to the collectors.

08:39.320 --> 08:45.560  
Next, the Custom Rule Engine (CRE) tests every event or flow against all enabled rules.

08:45.960 --> 08:48.800  
Matched rules might trigger the creation of an offense

08:49.280 --> 08:53.760  
or create a CRE event  
that then triggers the creation of an offense.

08:54.560 --> 08:59.960  
Multiple matched events, flows, and matched rules  
might correlate into a single offense,

09:00.000 --> 09:05.080  
and a single event or flow  
can be correlated into multiple offenses as well.

09:05.600 --> 09:12.880  
By default, rules are tested against events or flows  
that are received by a single Event Processor (local rules).

09:13.200 --> 09:16.960  
However, Global Cross-Correlation (GCC)  
allows rules testing across

09:17.000 --> 09:20.880

multiple Event Processors  
in the QRadar SIEM deployment.

09:21.920 --> 09:26.320  
From the CRE, the Exit Filter sends any events or flows  
that have been marked

09:26.360 --> 09:29.960  
for further processing  
by the Magistrate component on the Console.

09:30.920 --> 09:34.520  
Every event and flow is  
then sent to the Event Storage Filter,

09:34.720 --> 09:38.160  
where they are stored in the events or flows  
Ariel database.

09:39.040 --> 09:46.960  
If a new port or host is detected at this time, an asset profile needs to  
be updated or created in the PostgreSQL database.

09:47.760 --> 09:53.040  
The Host Profiler, also known as the Host  
Profiling Filter, sends the collected information

09:53.080 --> 09:57.320  
about the new host to the Console  
so that an asset can be created or updated.

09:58.640 --> 10:01.960  
Finally, if an analyst has defined any searches to collect

10:02.000 --> 10:06.480  
and investigate specific sets of data, events  
and flow records are accumulated

10:06.520 --> 10:10.400  
every minute and stored in the accumulations  
Ariel database tables.

10:10.960 --> 10:16.080  
These accumulations are used to create time-series  
statistical metadata for dashboards,

10:16.280 --> 10:22.200  
event and flow forensics and searching, reporting,  
and the Anomaly Detection Engine on the Console.

10:22.520 --> 10:27.240  
Accumulated time intervals can be defined  
as 1 minute, 1 hour, and 1 day.

10:28.520 --> 10:32.720  
The Accumulator is a distributed component  
that operates on each Event Processor.

10:35.480 --> 10:41.560  
The Console receives data from the deployed Event Processors for further  
analysis by the Magistrate component,

10:41.560 --> 10:45.960  
which creates, manages, and stores offenses  
in the PostgreSQL database.

10:47.320 --> 10:51.280  
These offenses are then brought to the analyst's  
attention in the user interface.

10:51.960 --> 10:55.600  
The Magistrate instructs  
the Ariel Proxy Server to gather information

10:55.640 --> 10:59.600  
about all related events and flows  
that triggered the creation of an offense.

11:00.320 --> 11:04.360  
The collected data is then available  
for further investigation by the analyst.

11:05.680 --> 11:11.160  
If data is collected from multiple Event Processors,  
the Console's Custom Rules Engine can use

11:11.200 --> 11:16.440  
Global Cross-Correlation to test rules and data  
from all deployed Event Processors.

11:17.880 --> 11:22.080  
This helps to locate more complex attacks,  
which can span across the overall

11:22.120 --> 11:26.960  
IT infrastructure and are not confined  
to being detected by a single Event Processor.

11:28.160 --> 11:32.240  
The Vulnerability Information Server detects  
and creates new assets

11:32.360 --> 11:36.280  
or adds information,  
such as open ports or discovered services

11:36.440 --> 11:41.840  
to existing assets, based on information

from the Host Profiler on the Event Processors.

11:42.520 --> 11:48.800

This happens when hosts, services, or vulnerabilities that cannot be mapped to existing assets are discovered.

11:49.360 --> 11:53.480

The Vulnerability Information Server automatically checks the asset information

11:53.520 --> 11:57.920

against uploaded vulnerability information by using flow information.

11:58.560 --> 12:05.640

The Anomaly Detection Engine searches the Accumulator databases for anomalies which are then used for offense evaluation.

12:06.680 --> 12:09.800

There are three categories of Anomaly Detection Rule types.

12:11.720 --> 12:18.000

The Threshold rule examines and numeric range, such as greater than, less than, or a particular range.

12:18.800 --> 12:22.200

This rule can help detect the bandwidth of an application,

12:22.200 --> 12:27.840

the number of users connected to a VPN, or a large and unusual outbound data transfer.

12:29.080 --> 12:34.560

The Anomaly rule looks at a change in short term when comparing against a longer timeframe.

12:34.960 --> 12:40.400

This can help locate new service activity or a change in the bandwidth volume on a specific link.

12:41.240 --> 12:47.080

And the Behavioral rule can detect changes from the same time the day before or the previous week.

12:47.600 --> 12:54.560

This includes mail traffic. For example, the increase on external or SMTP server traffic, which might be a relay.

12:55.440 --> 13:02.360

This rule can also be used for regular IT services, such as a backup monitoring, where the rule would trigger if a backup failed.

13:04.240 --> 13:08.760

Let's now take a closer look at how offenses are being managed by the Magistrate component.

13:10.200 --> 13:17.560

As seen earlier, rules can correlate events and flows into a single offense, and a single event or flow can belong to multiple offenses.

13:18.800 --> 13:22.400

While rules are tested, they might lead to the creation of an offense.

13:23.440 --> 13:26.040

Pending offenses tag the events and flows

13:26.040 --> 13:30.680

while the rule that triggered the creation of the offense remains at least partially matched.

13:31.320 --> 13:34.440

A maximum of 100,000 offenses can be stored.

13:35.280 --> 13:41.240

Rules are stored in the PostgreSQL database and replicated to Event and Flow Processors in your deployment.

13:42.400 --> 13:49.640

Events And flows that are tagged by the Custom Rules Engine for further processing are handed over to the Console through the Exit Filter.

13:52.320 --> 13:56.600

Now that you know the details of every component in a QRadar SIEM architecture,

13:56.720 --> 14:00.120

let's recap this section by examining a captured event

14:00.120 --> 14:02.960

from the time when it arrives at the Event Collector

14:02.960 --> 14:08.440

and follow it as it proceeds through correlation, accumulation and storage on the Event Processor

14:08.440 --> 14:12.040

until it ends up as part of a larger offense on the Console.